

A Theoretical Guide on Differentiation and a Practical Guide to JAX

(or, how never to compute a derivative by hand again)

May 19, 2022

The Basics

JAX is a lot like numpy

```
import numpy as np
import jax.numpy as jnp

x = np.array([1., 2., 3.])
y = np.array([0., 1., 5.])
z = np.array([2., 3., 4.])
print(f"with numpy: {x * y + z}")

x_ = jnp.array([1., 2., 3.])
y_ = jnp.array([0., 1., 5.])
z_ = jnp.array([2., 3., 4.])
print(f"with jax: {x_ * y_ + z_}")
```

✓ 0.5s

Python

JAX is a lot like numpy

```
import numpy as np
import jax.numpy as jnp

x = np.array([1., 2., 3.])
y = np.array([0., 1., 5.])
z = np.array([2., 3., 4.])
print(f"with numpy: {x * y + z}")

x_ = jnp.array([1., 2., 3.])
y_ = jnp.array([0., 1., 5.])
z_ = jnp.array([2., 3., 4.])
print(f"with jax: {x_ * y_ + z_}")
```

✓ 0.5s

Python

with numpy: [2. 5. 19.]

with jax: [2. 5. 19.]

Most np functions have a jnp counterpart

```
np_linspace = np.linspace(0, 1, 3)
jnp_linspace = jnp.linspace(0, 1, 3)

print(f"np_linspace: {np_linspace}")
print(f"jnp_linspace: {jnp_linspace}\n")

A = np.random.normal(size = (10, 10))
np_2norm = np.linalg.norm(A, ord = 2)
jnp_2norm = jnp.linalg.norm(A, ord = 2) # notice that we can feed np array into jnp function!

print(f"np_2norm: {np_2norm:.4f}")
print(f"jnp_2norm: {jnp_2norm:.4f}\n")
```

✓ 0.6s

Python

Most np functions have a jnp counterpart

```
np_linspace = np.linspace(0, 1, 3)
jnp_linspace = jnp.linspace(0, 1, 3)

print(f"np_linspace: {np_linspace}")
print(f"jnp_linspace: {jnp_linspace}\n")

A = np.random.normal(size = (10, 10))
np_2norm = np.linalg.norm(A, ord = 2)
jnp_2norm = jnp.linalg.norm(A, ord = 2) # notice that we can feed np array into jnp function!

print(f"np_2norm: {np_2norm:.4f}")
print(f"jnp_2norm: {jnp_2norm:.4f}\n")
```

✓ 0.6s

Python

```
np_linspace: [0.  0.5 1. ]
jnp_linspace: [0.  0.5 1. ]
```

```
np_2norm: 5.2232
jnp_2norm: 5.2232
```

matplotlib can accept jnp arrays

```
import matplotlib.pyplot as plt
```

```
x = jnp.linspace(0, 2 * jnp.pi)
```

```
y = jnp.sin(x)
```

```
plt.plot(x, y)
```

✓ 0.9s

Python

matplotlib can accept jnp arrays

```
import matplotlib.pyplot as plt
```

```
x = jnp.linspace(0, 2 * jnp.pi)
```

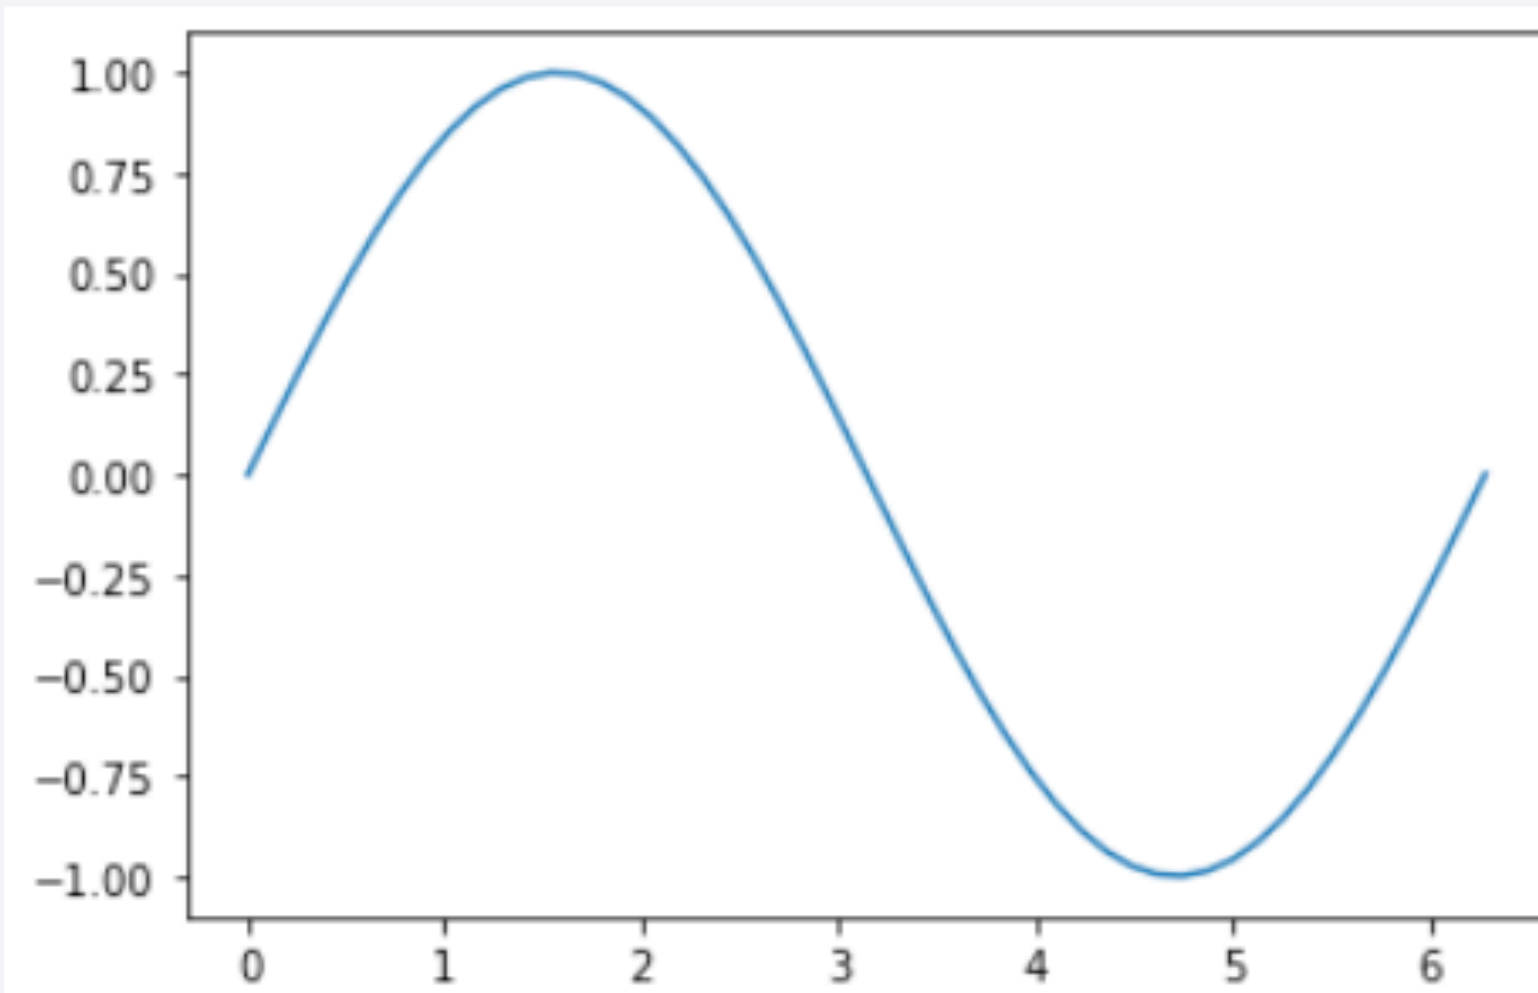
```
y = jnp.sin(x)
```

```
plt.plot(x, y)
```

✓ 0.9s

Python

[<matplotlib.lines.Line2D at 0x7f5d24561dc0>]



RNG is different

```
from jax import random

key = random.PRNGKey(0) # initialize key from your favorite integer

A = random.normal(key, shape = (2, 2))
print(f"A: {A}")
B = random.normal(key, shape = (2, 2)) # will get the same thing as A!
print(f"B: {B}")

key, subkey = random.split(key) # new key and subkey are different from original key
C = random.normal(key, shape = (2, 2))
print(f"C: {C}")
D = random.normal(subkey, shape = (2, 2))
print(f"D: {D}")
```

✓ 0.2s

Python

RNG is different

```
from jax import random

key = random.PRNGKey(0) # initialize key from your favorite integer

A = random.normal(key, shape = (2, 2))
print(f"A: {A}")
B = random.normal(key, shape = (2, 2)) # will get the same thing as A!
print(f"B: {B}")

key, subkey = random.split(key) # new key and subkey are different from original key
C = random.normal(key, shape = (2, 2))
print(f"C: {C}")
D = random.normal(subkey, shape = (2, 2))
print(f"D: {D}")
```

✓ 0.2s

Python

```
A: [[ 1.8160863 -0.75488514]
     [ 0.33988908 -0.53483534]]
B: [[ 1.8160863 -0.75488514]
     [ 0.33988908 -0.53483534]]
C: [[ 0.13893168  1.370668 ]
     [-0.53116107  0.02033782]]
D: [[ 1.1378784 -0.14331433]
     [-0.59153634  0.79466224]]
```

Updating arrays is different

in numpy, we can do in-place array updates:

```
A = np.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0
print(f"this is the new A in numpy: {A}")
```

✓ 0.3s

Python

in jax, this doesn't work:

```
A = jnp.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0 # this will throw an error
```

✗ 1.2s

Python

instead, we have to do this:

```
A = A.at[0].set(0)
print(f"this is the new A: {A}")
```

✓ 0.1s

Python

Updating arrays is different

in numpy, we can do in-place array updates:

```
A = np.array([1, 2, 3])  
print(f"this is A: {A}")  
  
A[0] = 0  
print(f"this is the new A in numpy: {A}")
```

✓ 0.3s

Python

this is A: [1 2 3]

this is the new A in numpy: [0 2 3]

in jax, this doesn't work:

```
A = jnp.array([1, 2, 3])  
print(f"this is A: {A}")  
  
A[0] = 0 # this will throw an error
```

✗ 1.2s

Python

instead, we have to do this:

```
A = A.at[0].set(0)  
print(f"this is the new A: {A}")
```

✓ 0.1s

Python

Updating arrays is different

in numpy, we can do in-place array updates:

```
A = np.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0
print(f"this is the new A in numpy: {A}")
```

✓ 0.3s

Python

```
this is A: [1 2 3]
this is the new A in numpy: [0 2 3]
```

in jax, this doesn't work:

```
A = jnp.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0 # this will throw an error
```

✗ 1.2s

Python

```
this is A: [1 2 3]
```

```
-----
TypeError                                Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 15' in <cell line: 5>()
      2 A = jnp.array([1, 2, 3])
      3 print(f"this is A: {A}")
----> 5 A[0] = 0
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/_src/numpy/lax_numpy.py:4512, in _unimplemented_setitem(self, i, x)
    4507 def _unimplemented_setitem(self, i, x):
    4508     msg = ("'{}' object does not support item assignment. JAX arrays are "
    4509           "immutable. Instead of ``x[idx] = y``, use ``x = x.at[idx].set(y)`` "
    4510           "or another .at[] method: "
    4511           "https://jax.readthedocs.io/en/latest/_autosummary/jax.numpy.ndarray.at.html")
-> 4512     raise TypeError(msg.format(type(self)))
```

```
TypeError: '<class 'jaxlib.xla_extension.DeviceArray'>' object does not support item assignment. JAX arrays are immutable. Instead of ``x[id
x] = y``, use ``x = x.at[idx].set(y)`` or another .at[] method: https://jax.readthedocs.io/en/latest/_autosummary/jax.numpy.ndarray.at.html
```

instead, we have to do this:

```
A = A.at[0].set(0)
print(f"this is the new A: {A}")
```

✓ 0.1s

Python

Updating arrays is different

in numpy, we can do in-place array updates:

```
A = np.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0
print(f"this is the new A in numpy: {A}")
```

✓ 0.3s

Python

this is A: [1 2 3]

this is the new A in numpy: [0 2 3]

in jax, this doesn't work:

```
A = jnp.array([1, 2, 3])
print(f"this is A: {A}")

A[0] = 0 # this will throw an error
```

✗ 1.2s

Python

this is A: [1 2 3]

```
-----
TypeError                                Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 15' in <cell line: 5>()
      2 A = jnp.array([1, 2, 3])
      3 print(f"this is A: {A}")
----> 5 A[0] = 0
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/_src/numpy/lax_numpy.py:4512, in _unimplemented_setitem(self, i, x)
    4507 def _unimplemented_setitem(self, i, x):
    4508     msg = ("'{}' object does not support item assignment. JAX arrays are "
    4509           "immutable. Instead of ``x[idx] = y``, use ``x = x.at[idx].set(y)`` "
    4510           "or another .at[] method: "
    4511           "https://jax.readthedocs.io/en/latest/_autosummary/jax.numpy.ndarray.at.html")
-> 4512     raise TypeError(msg.format(type(self)))
```

```
TypeError: '<class 'jaxlib.xla_extension.DeviceArray'>' object does not support item assignment. JAX arrays are immutable. Instead of ``x[id
x] = y``, use ``x = x.at[idx].set(y)`` or another .at[] method: https://jax.readthedocs.io/en/latest/_autosummary/jax.numpy.ndarray.at.html
```

instead, we have to do this:

```
A = A.at[0].set(0)
print(f"this is the new A: {A}")
```

✓ 0.1s

Python

this is the new A: [0 2 3]

Computing Gradients

JAX is highly functional

Use `grad()` to map a scalar-valued function to its derivative

```
from jax import grad
```

```
sin_grad = grad(jnp.sin)
```

```
print(f"sin'(1): {sin_grad(1.0):.4f}")
```

```
print(f"cos(1): {jnp.cos(1.0):.4f}")
```

✓ 0.4s

Python

JAX is highly functional

Use `grad()` to map a scalar-valued function to its derivative

```
from jax import grad
```

```
sin_grad = grad(jnp.sin)
```

```
print(f"sin'(1): {sin_grad(1.0):.4f}")
```

```
print(f"cos(1): {jnp.cos(1.0):.4f}")
```

✓ 0.4s

Python

```
sin'(1): 0.5403
```

```
cos(1): 0.5403
```

Get fancy with function composition

Let $f(x) = \sin(x^2) + 3x^2$, then $f'(x) = 2x \cos(x^2) + 6x$

```
def f(x):  
    return jnp.sin(x**2) + 3 * x**2  
  
def grad_f_manual(x):  
    return 2 * x * jnp.cos(x**2) + 6 * x  
  
grad_f_auto = grad(f)  
  
x = 2.0  
print(f"manual f'(1): {grad_f_manual(x):.4f}")  
print(f"auto f'(1): {grad_f_auto(x):.4f}")
```

✓ 0.1s

Python

Get fancy with function composition

Let $f(x) = \sin(x^2) + 3x^2$, then $f'(x) = 2x \cos(x^2) + 6x$

```
def f(x):  
    return jnp.sin(x**2) + 3 * x**2  
  
def grad_f_manual(x):  
    return 2 * x * jnp.cos(x**2) + 6 * x  
  
grad_f_auto = grad(f)  
  
x = 2.0  
print(f"manual f'(1): {grad_f_manual(x):.4f}")  
print(f"auto f'(1): {grad_f_auto(x):.4f}")
```

✓ 0.1s

Python

manual f'(1): 9.3854

auto f'(1): 9.3854

Multivariate Functions

$$f(x, y, z) = x^2 + 3 \sin(y) + \cos(z)$$

$$\nabla f(x, y, z) = (2x, 3 \cos(y), -\sin(z))$$

```
# denote X = (x, y, z)

def f(X):
    return X[0]**2 + 3 * jnp.sin(X[1]) + jnp.cos(X[2])

def grad_f_manual(X):
    return jnp.array([2 * X[0], 3 * jnp.cos(X[1]), -jnp.sin(X[2])])

grad_f_auto = grad(f)

X = jnp.array([2.0, 3.0, 1.0])
print(f"manual f'(1): {grad_f_manual(X)}")
print(f"auto f'(1): {grad_f_auto(X)}")
```

✓ 0.5s

Python

Multivariate Functions

$$f(x, y, z) = x^2 + 3 \sin(y) + \cos(z)$$

$$\nabla f(x, y, z) = (2x, 3 \cos(y), -\sin(z))$$

```
# denote X = (x, y, z)

def f(X):
    return X[0]**2 + 3 * jnp.sin(X[1]) + jnp.cos(X[2])

def grad_f_manual(X):
    return jnp.array([2 * X[0], 3 * jnp.cos(X[1]), -jnp.sin(X[2])])

grad_f_auto = grad(f)

X = jnp.array([2.0, 3.0, 1.0])
print(f"manual f'(1): {grad_f_manual(X)}")
print(f"auto f'(1): {grad_f_auto(X)}")
```

✓ 0.5s Python

```
manual f'(1): [ 4.          -2.9699774 -0.84147096]
auto f'(1): [ 4.          -2.9699774 -0.84147096]
```

Differentiate with respect to dictionaries

$$f(x, y, z) = x^2 + 3 \sin(y) + \cos(z)$$

$$\nabla f(x, y, z) = (2x, 3 \cos(y), -\sin(z))$$

```
def f(X):  
    return X['x']**2 + 3 * jnp.sin(X['y']) + jnp.cos(X['z'])  
  
def grad_f_manual(X):  
    return jnp.array([2 * X['x'], 3 * jnp.cos(X['y']), -jnp.sin(X['z'])])  
  
grad_f_auto = grad(f)  
  
X = {'x' : 2.0, 'y' : 3.0, 'z' : 1.0}  
print(f"manual f'(1): {grad_f_manual(X)}")  
print(f"auto f'(1): {grad_f_auto(X)}")
```

✓ 0.5s

Python

Differentiate with respect to dictionaries

$$f(x, y, z) = x^2 + 3 \sin(y) + \cos(z)$$

$$\nabla f(x, y, z) = (2x, 3 \cos(y), -\sin(z))$$

```
def f(X):  
    return X['x']**2 + 3 * jnp.sin(X['y']) + jnp.cos(X['z'])  
  
def grad_f_manual(X):  
    return jnp.array([2 * X['x'], 3 * jnp.cos(X['y']), -jnp.sin(X['z'])])  
  
grad_f_auto = grad(f)  
  
X = {'x' : 2.0, 'y' : 3.0, 'z' : 1.0}  
print(f"manual f'(1): {grad_f_manual(X)}")  
print(f"auto f'(1): {grad_f_auto(X)}")
```

✓ 0.5s

Python

```
manual f'(1): [ 4.          -2.9699774 -0.84147096]  
auto f'(1): {'x': DeviceArray(4., dtype=float32, weak_type=True), 'y':  
DeviceArray(-2.9699774, dtype=float32, weak_type=True), 'z':  
DeviceArray(-0.84147096, dtype=float32, weak_type=True)}
```

Multiple multivariate arguments

$$f(x, y) = \|x\|^2 + \|y\| + \langle x, y \rangle$$

$$D_x f(x, y) = 2x + y$$

$$D_y f(x, y) = \frac{y}{\|y\|} + x$$

$$x, y \in \mathbb{R}^3$$

```
def f(x, y):  
    return jnp.linalg.norm(x)**2 + jnp.linalg.norm(y) + jnp.inner(x, y)  
  
Dxf = grad(f, 0) # take the derivative wrt the zeroth argument  
Dyf = grad(f, 1) # take the derivative wrt the first argument  
  
x = jnp.array([1., 2., 3.])  
y = jnp.array([0, -1., 1.])  
  
print(f"derivative wrt x: {Dxf(x, y)}")  
print(f"manually computed: {2 * x + y}\n")  
print(f"derivative wrt y: {Dyf(x, y)}")  
print(f"manually computed: {y / jnp.linalg.norm(y) + x}")
```

✓ 0.8s

Python

Multiple multivariate arguments

$$f(x, y) = \|x\|^2 + \|y\| + \langle x, y \rangle$$

$$D_x f(x, y) = 2x + y$$

$$D_y f(x, y) = \frac{y}{\|y\|} + x$$

$$x, y \in \mathbb{R}^3$$

```
def f(x, y):  
    return jnp.linalg.norm(x)**2 + jnp.linalg.norm(y) + jnp.inner(x, y)  
  
Dxf = grad(f, 0) # take the derivative wrt the zeroth argument  
Dyf = grad(f, 1) # take the derivative wrt the first argument  
  
x = jnp.array([1., 2., 3.])  
y = jnp.array([0, -1., 1.])  
  
print(f"derivative wrt x: {Dxf(x, y)}")  
print(f"manually computed: {2 * x + y}\n")  
print(f"derivative wrt y: {Dyf(x, y)}")  
print(f"manually computed: {y / jnp.linalg.norm(y) + x}")
```

✓ 0.8s

Python

```
derivative wrt x: [2. 3. 7.]  
manually computed: [2. 3. 7.]
```

```
derivative wrt y: [1.          1.2928932 3.7071068]  
manually computed: [1.          1.2928932 3.7071068]
```

value_and_grad()

Useful for first-order optimization methods

```
from jax import value_and_grad

def f(x):
    return x ** 2

value, grad = value_and_grad(f)(3.0)
print(f"value: {value}")
print(f"grad: {grad}")
```

✓ 0.4s

Python

value_and_grad()

Useful for first-order optimization methods

```
from jax import value_and_grad
```

```
def f(x):  
    return x ** 2
```

```
value, grad = value_and_grad(f)(3.0)  
print(f"value: {value}")  
print(f"grad: {grad}")
```

✓ 0.4s

Python

value: 9.0

grad: 6.0

What is Automatic Differentiation?

Not finite differences or symbolic differentiation

Not finite differences or symbolic differentiation

- Finite differences: $\frac{d}{dx}f(x) \approx \frac{f(x+h) - f(x)}{h}$

Not finite differences or symbolic differentiation

- Finite differences: $\frac{d}{dx}f(x) \approx \frac{f(x+h) - f(x)}{h}$
 - Numerically unstable

Not finite differences or symbolic differentiation

- Finite differences: $\frac{d}{dx}f(x) \approx \frac{f(x+h) - f(x)}{h}$
 - Numerically unstable
 - Expensive

Not finite differences or symbolic differentiation

- Finite differences: $\frac{d}{dx}f(x) \approx \frac{f(x+h) - f(x)}{h}$
 - Numerically unstable
 - Expensive
- Symbolic differentiation:

Not finite differences or symbolic differentiation

- Finite differences: $\frac{d}{dx}f(x) \approx \frac{f(x+h) - f(x)}{h}$
 - Numerically unstable
 - Expensive
- Symbolic differentiation:

```
In[1]:= D[Log[1 + Exp[w2 * Log[1 + Exp[w1 * x + b1]]] + b2]], w1]
```

Out[1]=

$$\frac{e^{b_1+b_2+w_1 x+w_2 \operatorname{Log}\left[1+e^{b_1+w_1 x}\right]} w_2 x}{\left(1+e^{b_1+w_1 x}\right) \times \left(1+e^{b_2+w_2 \operatorname{Log}\left[1+e^{b_1+w_1 x}\right]}\right)}$$

Chain rule

Chain rule

$$h(x) = f(g(x))$$

$$\frac{dh}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

Chain rule

$$h(x) = f(g(x))$$

$$\frac{dh}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

$$\frac{d}{dx} f(g_1(x), \dots, g_K(x)) = \sum_{k=1}^K \left(\frac{d}{dx} g_k(x) \right) D_k f(g_1(x), \dots, g_K(x))$$

Chain rule

$$h(x) = f(g(x))$$

$$\frac{dh}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

$$\frac{d}{dx} f(g_1(x), \dots, g_K(x)) = \sum_{k=1}^K \left(\frac{d}{dx} g_k(x) \right) D_k f(g_1(x), \dots, g_K(x))$$

Idea: the chain rule is modular; as long as we know the derivatives of elementary functions, we can compute the derivatives of arbitrary compositions of these functions

Example: regularized univariate linear predictor

$$z = mx + b$$

$$\hat{y} = \sigma(z)$$


$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$


Example: regularized univariate linear predictor

$$\begin{aligned} z &= mx + b \\ \hat{y} &= \sigma(z) \\ L^{m,b}(\hat{y}, y) &= \frac{1}{2}(\hat{y} - y)^2 \\ R(m) &= \frac{1}{2}m^2 \\ L_R^{m,b} &= L^{m,b}(\hat{y}, y) + \lambda R(m) \end{aligned}$$



Example: regularized univariate linear predictor

$$\begin{aligned} z &= mx + b \\ \hat{y} &= \sigma(z) \\ L^{m,b}(\hat{y}, y) &= \frac{1}{2}(\hat{y} - y)^2 \\ R(m) &= \frac{1}{2}m^2 \\ L_R^{m,b} &= L^{m,b}(\hat{y}, y) + \lambda R(m) \end{aligned}$$



parent \rightarrow child, derivative w.r.t.
parent requires derivatives w.r.t
children

Example: regularized univariate linear predictor


$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$




parent \rightarrow child, derivative w.r.t.
parent requires derivatives w.r.t
children

Automatic differentiation algorithm:
work backwards through the
computation tree, computing
derivatives of parents w.r.t children

Example: regularized univariate linear predictor

$$\begin{aligned}z &= mx + b \\ \hat{y} &= \sigma(z) \\ L^{m,b}(\hat{y}, y) &= \frac{1}{2}(\hat{y} - y)^2 \\ R(m) &= \frac{1}{2}m^2 \\ L_R^{m,b} &= L^{m,b}(\hat{y}, y) + \lambda R(m)\end{aligned}$$

$$D_{L_R} = 1$$



parent \rightarrow child, derivative w.r.t.
parent requires derivatives w.r.t
children

Automatic differentiation algorithm:
work backwards through the
computation tree, computing
derivatives of parents w.r.t children

Example: regularized univariate linear predictor


$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t.
parent requires derivatives w.r.t
children

Automatic differentiation algorithm:
 work backwards through the
 computation tree, computing
 derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

Example: regularized univariate linear predictor


$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t.
parent requires derivatives w.r.t
children

Automatic differentiation algorithm:
 work backwards through the
 computation tree, computing
 derivatives of parents w.r.t children


$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

Example: regularized univariate linear predictor

$$\begin{aligned}
 z &= mx + b \\
 \hat{y} &= \sigma(z) \\
 L^{m,b}(\hat{y}, y) &= \frac{1}{2}(\hat{y} - y)^2 \\
 R(m) &= \frac{1}{2}m^2 \\
 L_R^{m,b} &= L^{m,b}(\hat{y}, y) + \lambda R(m)
 \end{aligned}$$



parent → child, derivative w.r.t. parent requires derivatives w.r.t. children

Automatic differentiation algorithm:
work backwards through the
computation tree, computing
derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

$$D_{\hat{y}} = D_L \frac{dL}{dy} = (\hat{y} - y) D_L$$

Example: regularized univariate linear predictor


$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t. parent requires derivatives w.r.t. children

Automatic differentiation algorithm:
work backwards through the
computation tree, computing
derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

$$D_{\hat{y}} = D_L \frac{dL}{d\hat{y}} = (\hat{y} - y) D_L$$

$$D_z = D_{\hat{y}} \frac{d\hat{y}}{dz} = D_{\hat{y}} \sigma'(z)$$

Example: regularized univariate linear predictor


$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t. parent requires derivatives w.r.t. children

Automatic differentiation algorithm: work backwards through the computation tree, computing derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

$$D_{\hat{y}} = D_L \frac{dL}{d\hat{y}} = (\hat{y} - y) D_L$$

$$D_z = D_{\hat{y}} \frac{d\hat{y}}{dz} = D_{\hat{y}} \sigma'(z)$$

$$D_m = D_z \frac{dz}{dm} + D_R \frac{dR}{dm}$$

$$= D_z x + D_R m$$

Example: regularized univariate linear predictor

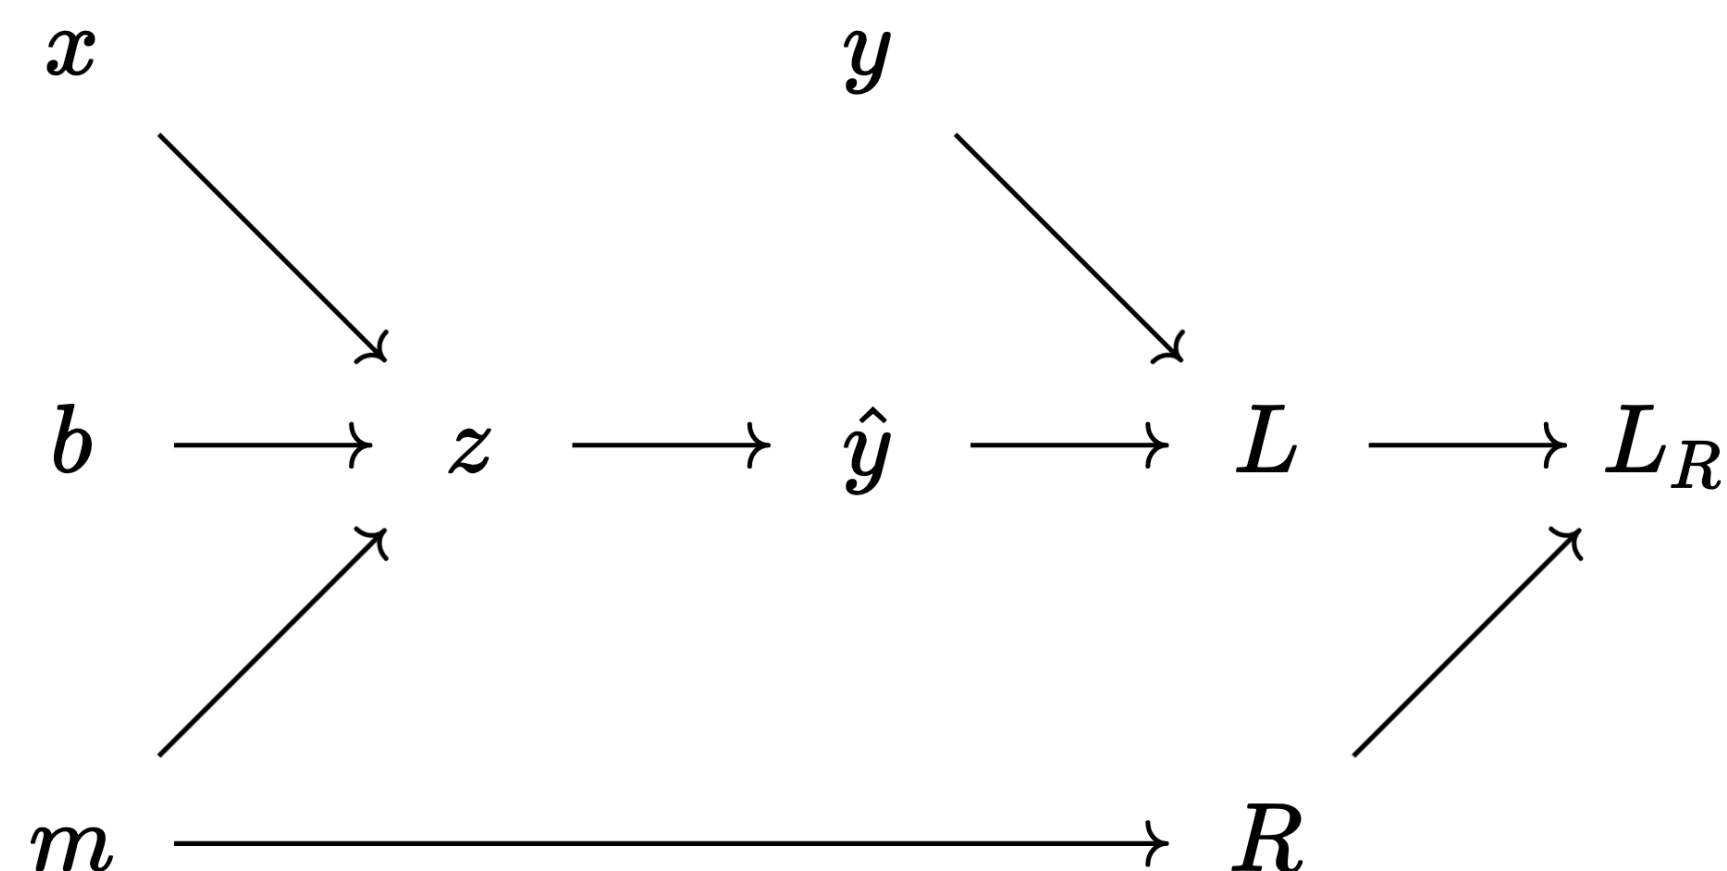
$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t. parent requires derivatives w.r.t. children

Automatic differentiation algorithm: work backwards through the computation tree, computing derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

$$D_{\hat{y}} = D_L \frac{dL}{d\hat{y}} = (\hat{y} - y) D_L$$

$$D_z = D_{\hat{y}} \frac{d\hat{y}}{dz} = D_{\hat{y}} \sigma'(z)$$

$$D_m = D_z \frac{dz}{dm} + D_R \frac{dR}{dm}$$

$$= D_z x + D_R m$$

etc...

Example: regularized univariate linear predictor

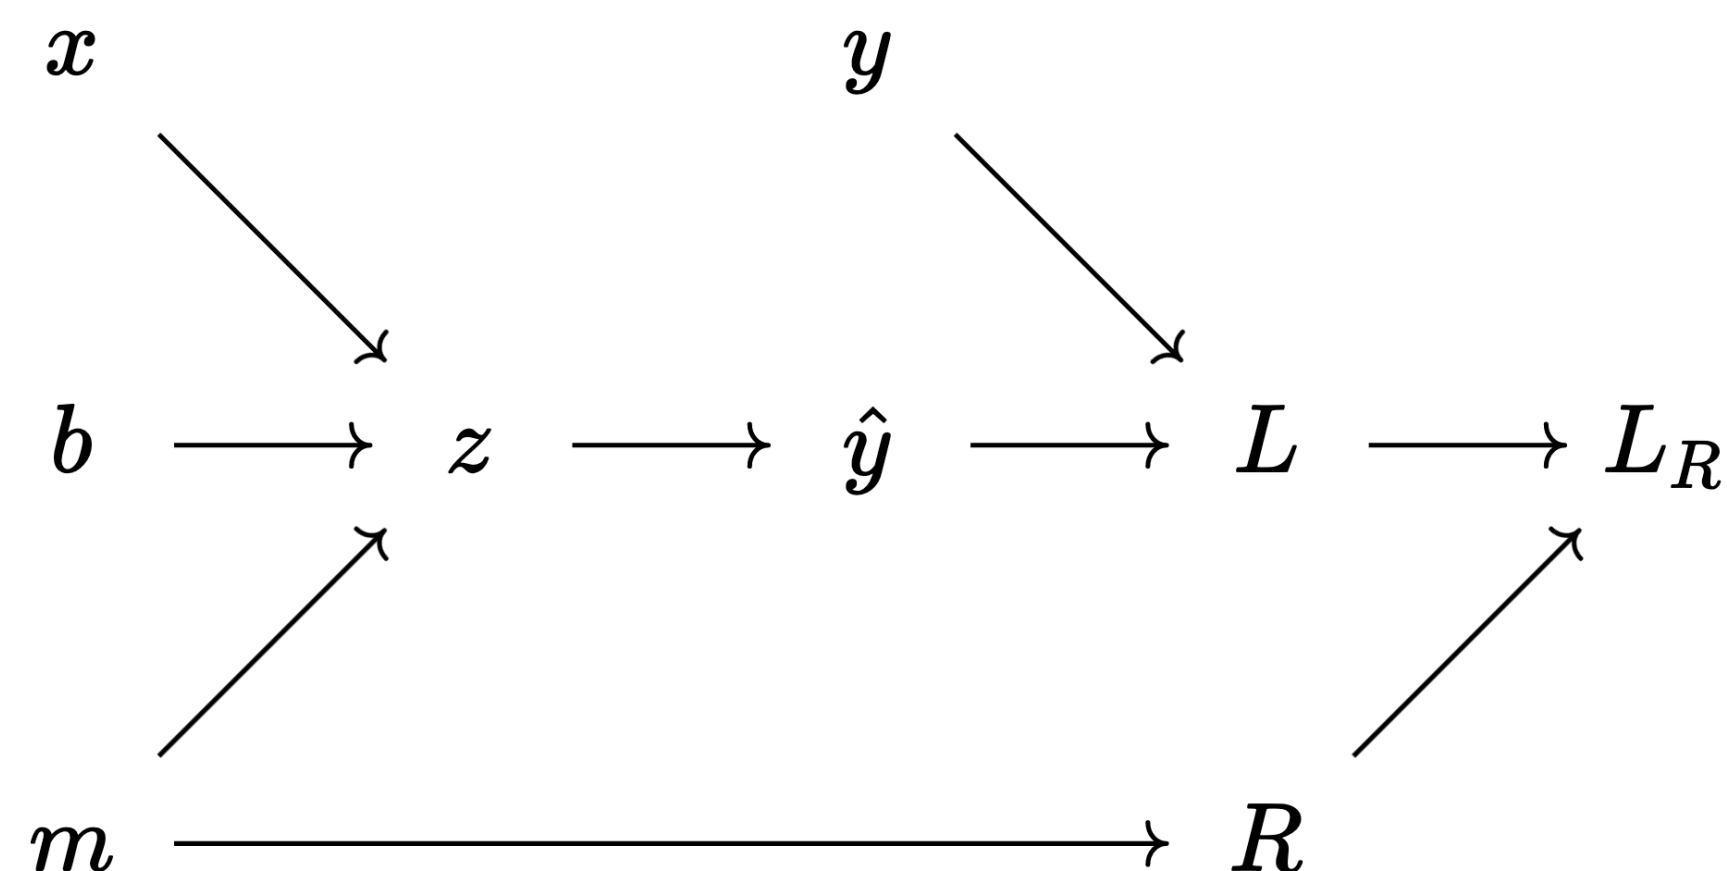
$$z = mx + b$$

$$\hat{y} = \sigma(z)$$

$$L^{m,b}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

$$R(m) = \frac{1}{2}m^2$$

$$L_R^{m,b} = L^{m,b}(\hat{y}, y) + \lambda R(m)$$



parent → child, derivative w.r.t. parent requires derivatives w.r.t. children

Automatic differentiation algorithm: work backwards through the computation tree, computing derivatives of parents w.r.t children

$$D_{L_R} = 1$$

$$D_L = D_{L_R} \frac{dL_R}{dL} = D_{L_R}$$

$$D_R = D_{L_R} \frac{dL_R}{dR} = \lambda D_{L_R}$$

$$D_{\hat{y}} = D_L \frac{dL}{d\hat{y}} = (\hat{y} - y) D_L$$

$$D_z = D_{\hat{y}} \frac{d\hat{y}}{dz} = D_{\hat{y}} \sigma'(z)$$

$$D_m = D_z \frac{dz}{dm} + D_R \frac{dR}{dm}$$

$$= D_z x + D_R m$$

etc...

Key insight: modular! No need to recompute things!

Jacobians

What is a Jacobian?

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- What is the Jacobian of f ?

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- What is the Jacobian of f ?
- Calculus student: matrix of partial derivatives

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- What is the Jacobian of f ?
- Calculus student: matrix of partial derivatives

$$Jf(x) = \left[\frac{\partial f}{\partial x_{ij}} \right]_{ij} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_j} \\ \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_j} \end{bmatrix}$$

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- What is the Jacobian of f ?
- Calculus student: matrix of partial derivatives

$$Jf(x) = \left[\frac{\partial f}{\partial x_{ij}} \right]_{ij} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_j} \\ \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_j} \end{bmatrix}$$

- Functional analyst: linear operator from $\mathbb{R}^n \rightarrow \mathbb{R}^m$

What is a Jacobian?

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- What is the Jacobian of f ?
- Calculus student: matrix of partial derivatives

$$Jf(x) = \left[\frac{\partial f}{\partial x_{ij}} \right]_{ij} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_j} \\ \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_j} \end{bmatrix}$$

- Functional analyst: linear operator from $\mathbb{R}^n \rightarrow \mathbb{R}^m$
- Differential geometer: a map from the tangent bundle of \mathbb{R}^n to \mathbb{R}^m

Thinking about Jacobians like a Functional Analyst

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

$$Jf(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$$
$$v \mapsto Jf(x)v$$

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

$$Jf(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$v \mapsto Jf(x)v$$

- The domain of $Jf(x)$ is the tangent space to the domain of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, denoted $T_x\mathbb{R}^n$; this just happens to be a *copy* of \mathbb{R}^n

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

$$Jf(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$v \mapsto Jf(x)v$$

- The domain of $Jf(x)$ is the tangent space to the domain of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, denoted $T_x\mathbb{R}^n$; this just happens to be a *copy* of \mathbb{R}^n
 - Tangent space to a manifold is the vector space best approximating that manifold

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

$$Jf(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$v \mapsto Jf(x)v$$

- The domain of $Jf(x)$ is the tangent space to the domain of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, denoted $T_x\mathbb{R}^n$; this just happens to be a *copy* of \mathbb{R}^n
 - Tangent space to a manifold is the vector space best approximating that manifold
 - Tangent vectors v and domain vectors x are different types of objects

Thinking about Jacobians like a Functional Analyst

- Linear operator vs. the matrix *representing* that operator!
- Jacobian as a linear operator:

$$Jf(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$v \mapsto Jf(x)v$$

- The domain of $Jf(x)$ is the tangent space to the domain of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, denoted $T_x \mathbb{R}^n$; this just happens to be a *copy* of \mathbb{R}^n
 - Tangent space to a manifold is the vector space best approximating that manifold
 - Tangent vectors v and domain vectors x are different types of objects
- Similarly, the range of $Jf(x)$ is the tangent space to the range of f at $f(x)$, denoted $T_{f(x)} \mathbb{R}^m$

Thinking about Jacobians like a Differential Geometer

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

$$Jf : \mathbb{R}^n \rightarrow \underbrace{(\mathbb{R}^n \rightarrow \mathbb{R}^m)}_{\text{an operator}}$$

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

$$Jf : \mathbb{R}^n \rightarrow \underbrace{(\mathbb{R}^n \rightarrow \mathbb{R}^m)}_{\text{an operator}}$$

- Better yet, we can think of Jf as taking in two inputs; domain is *tangent bundle* of \mathbb{R}^n :

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

$$Jf : \mathbb{R}^n \rightarrow \underbrace{(\mathbb{R}^n \rightarrow \mathbb{R}^m)}_{\text{an operator}}$$

- Better yet, we can think of Jf as taking in two inputs; domain is *tangent bundle* of \mathbb{R}^n :

$$Jf : (x, v) \mapsto Jf(x)v$$

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

$$Jf : \mathbb{R}^n \rightarrow \underbrace{(\mathbb{R}^n \rightarrow \mathbb{R}^m)}_{\text{an operator}}$$

- Better yet, we can think of Jf as taking in two inputs; domain is *tangent bundle* of \mathbb{R}^n :

$$Jf : (x, v) \mapsto Jf(x)v$$

- When thinking of Jacobians in this way: we call this the Jacobian-vector product (JVP)

Thinking about Jacobians like a Differential Geometer

- If we don't fix x , then Jf is a function that first takes a domain vector $x \in \mathbb{R}^n$, which then defines a linear operator at that point:

$$Jf : \mathbb{R}^n \rightarrow \underbrace{(\mathbb{R}^n \rightarrow \mathbb{R}^m)}_{\text{an operator}}$$

- Better yet, we can think of Jf as taking in two inputs; domain is *tangent bundle* of \mathbb{R}^n :

$$Jf : (x, v) \mapsto Jf(x)v$$

- When thinking of Jacobians in this way: we call this the Jacobian-vector product (JVP)

JAX's `jvp()`

$$f(x, y) = (x + 2y, \sin(x)e^y, y^3), f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$Jf(x, y) = \begin{bmatrix} 1 & 2 \\ \cos(x)e^y & \sin(x)e^y \\ 0 & 3y^2 \end{bmatrix} \in \mathbb{R}^{3 \times 2}$$

```
from jax import jvp

def f(x):
    return jnp.array([x[0] + 2 * x[1], jnp.sin(x[0]) * jnp.exp(x[1]), x[1]**3])

def Jf_manual(x):
    return jnp.array([
        [1, 2],
        [jnp.cos(x[0]) * jnp.exp(x[1]), jnp.sin(x[0]) * jnp.exp(x[1])],
        [0, 3 * x[1] **2]
    ])

# compute the jvp manually
def jvp_manual(x, v):
    return Jf_manual(x) @ v

x = jnp.array([1., 2.])
v = jnp.array([2., 3.])

print(f"manually computed jvp: {jvp_manual(x, v)}")

# compute the jvp using jax.jvp():

func_value, jvp_JAX = jvp(f, (x,), (v,)) # jvp returns both the function value and the jvp
print(f"JAX computed jvp: {jvp_JAX}")
```

✓ 0.3s

Python

JAX's `jvp()`

$$f(x, y) = (x + 2y, \sin(x)e^y, y^3), f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$Jf(x, y) = \begin{bmatrix} 1 & 2 \\ \cos(x)e^y & \sin(x)e^y \\ 0 & 3y^2 \end{bmatrix} \in \mathbb{R}^{3 \times 2}$$

```
from jax import jvp

def f(x):
    return jnp.array([x[0] + 2 * x[1], jnp.sin(x[0]) * jnp.exp(x[1]), x[1]**3])

def Jf_manual(x):
    return jnp.array([
        [1, 2],
        [jnp.cos(x[0]) * jnp.exp(x[1]), jnp.sin(x[0]) * jnp.exp(x[1])],
        [0, 3 * x[1] **2]
    ])

# compute the jvp manually
def jvp_manual(x, v):
    return Jf_manual(x) @ v

x = jnp.array([1., 2.])
v = jnp.array([2., 3.])

print(f"manually computed jvp: {jvp_manual(x, v)}")

# compute the jvp using jax.jvp():

func_value, jvp_JAX = jvp(f, (x,), (v,)) # jvp returns both the function value and the jvp
print(f"JAX computed jvp: {jvp_JAX}")
```

✓ 0.3s

Python

manually computed jvp: [8. 26.637676 36.]

JAX computed jvp: [8. 26.637676 36.]

But what if we want the full Jacobian matrix?

Post-multiplying a matrix with the j th standard basis vector reveals the j th column of the matrix, so can construct the full $m \times n$ Jacobian matrix with n JVPs:

```
def Jf_JVPs(x): # compute the full Jacobian at x using JVPs
    Jf = np.zeros(shape = (3, 2))
    for j in range(2):
        # compute the jth basis vector
        ej = np.zeros(shape = 2)
        ej[j] = 1

        # set the jth column equal to JVP with v = ej
        Jf[:, j] = jvp(f, (x,), (ej,))[1] # remember jvp returns function value and jvp

    return Jf

x = jnp.array([2., 3.])

print(f"Jf with JVPs: \n {Jf_JVPs(x)}")
print(f"Jf manual: \n {Jf_manual(x)}")
```

✓ 0.7s

Python

But what if we want the full Jacobian matrix?

Post-multiplying a matrix with the j th standard basis vector reveals the j th column of the matrix, so can construct the full $m \times n$ Jacobian matrix with n JVPs:

```
def Jf_JVPs(x): # compute the full Jacobian at x using JVPs
    Jf = np.zeros(shape = (3, 2))
    for j in range(2):
        # compute the jth basis vector
        ej = np.zeros(shape = 2)
        ej[j] = 1

        # set the jth column equal to JVP with v = ej
        Jf[:, j] = jvp(f, (x,), (ej,))[1] # remember jvp returns function value and jvp

    return Jf

x = jnp.array([2., 3.])

print(f"Jf with JVPs: \n {Jf_JVPs(x)}")
print(f"Jf manual: \n {Jf_manual(x)}")
```

✓ 0.7s

Python

Jf with JVPs:

```
[[ 1.         2.         ]
 [-8.35853291 18.26372719]
 [ 0.         27.         ]]
```

Jf manual:

```
[[ 1.         2.         ]
 [-8.358533   18.263727]
 [ 0.         27.         ]]
```

Forward-mode differentiation: `jacfwd()`

```
from jax import jacfwd
```

```
print(f"Jf with jacfwd(): \n {jacfwd(f)(x)}")
```

✓ 0.3s

Python

```
Jf with jacfwd():
```

```
[[ 1.          2.         ]
```

```
[-8.358533 18.263727]
```

```
[ 0.          27.         ]]
```

JVPs good for tall and skinny Jacobians

JVPs good for tall and skinny Jacobians

- JVPs compute $Jf(x)v$

JVPs good for tall and skinny Jacobians

- JVPs compute $Jf(x)v$
- By repeating this for $v = e_1, \dots, e_n$, we can reconstruct the full Jacobian matrix one column at a time; `jac_fwd()` does this for us under the hood

JVPs good for tall and skinny Jacobians

- JVPs compute $Jf(x)v$
- By repeating this for $v = e_1, \dots, e_n$, we can reconstruct the full Jacobian matrix one column at a time; `jac_fwd()` does this for us under the hood
- Better for tall and skinny Jacobians (a few variables mapping to a lot of variables)

JVPs good for tall and skinny Jacobians

- JVPs compute $Jf(x)v$
- By repeating this for $v = e_1, \dots, e_n$, we can reconstruct the full Jacobian matrix one column at a time; `jac_fwd()` does this for us under the hood
- Better for tall and skinny Jacobians (a few variables mapping to a lot of variables)
- But often more interested in short and fat Jacobians (a lot of variables mapping to a few variables), e.g., loss functions

JVPs good for tall and skinny Jacobians

- JVPs compute $Jf(x)v$
- By repeating this for $v = e_1, \dots, e_n$, we can reconstruct the full Jacobian matrix one column at a time; `jac_fwd()` does this for us under the hood
- Better for tall and skinny Jacobians (a few variables mapping to a lot of variables)
- But often more interested in short and fat Jacobians (a lot of variables mapping to a few variables), e.g., loss functions
- If only we could reconstruct a Jacobian matrix one row at a time...

Vector-Jacobian products (VJPs)

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix
- Vector Jacobian products: $(x, v^T) \mapsto v^T Jf(x)$

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix
- Vector Jacobian products: $(x, v^T) \mapsto v^T Jf(x)$
- v^T lives in the *cotangent space* (dual of the tangent space) of the range of $f(x)$

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix
- Vector Jacobian products: $(x, v^T) \mapsto v^T Jf(x)$
- v^T lives in the *cotangent space* (dual of the tangent space) of the range of $f(x)$
 - Remember $Jf(x)$ maps from the tangent space of \mathbb{R}^n at a point to the tangent space of \mathbb{R}^m at a point

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix
- Vector Jacobian products: $(x, v^T) \mapsto v^T Jf(x)$
- v^T lives in the *cotangent space* (dual of the tangent space) of the range of $f(x)$
 - Remember $Jf(x)$ maps from the tangent space of \mathbb{R}^n at a point to the tangent space of \mathbb{R}^m at a point
 - So v^T is a *linear functional* on the tangent space of \mathbb{R}^m : $v^T \in (T_{f(x)}\mathbb{R}^m)^* \cong \mathbb{R}^m$

Vector-Jacobian products (VJPs)

- If $Jf(x) \in \mathbb{R}^{m \times n}$ is a Jacobian matrix, then premultiplying with the basis vector $e_i^T \in \mathbb{R}^m$ will reveal the i th row of the Jacobian matrix
- Vector Jacobian products: $(x, v^T) \mapsto v^T Jf(x)$
- v^T lives in the *cotangent space* (dual of the tangent space) of the range of $f(x)$
 - Remember $Jf(x)$ maps from the tangent space of \mathbb{R}^n at a point to the tangent space of \mathbb{R}^m at a point
 - So v^T is a *linear functional* on the tangent space of \mathbb{R}^m : $v^T \in (T_{f(x)}\mathbb{R}^m)^* \cong \mathbb{R}^m$
- The n -dimensional vector $v^T Jf(x)$ is a row vector, making it a linear operator on \mathbb{R}^n , so it lives in the cotangent space of \mathbb{R}^n : $v^T Jf(x) \in (T_x\mathbb{R}^n)^* \cong \mathbb{R}^n$

$$(T_x \mathbb{R}^n)^* \xleftarrow{\text{vjp}} (T_{f(x)} \mathbb{R}^m)^*$$

$$T_x \mathbb{R}^n \xrightarrow{\text{jvp}} T_{f(x)} \mathbb{R}^m$$

$$\mathbb{R}^n \xrightarrow{f} \mathbb{R}^m$$

vjp() example

$$f(x, y) = (x + 2y, \sin(x)e^y, y^3), f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$Jf(x, y) = \begin{bmatrix} 1 & 2 \\ \cos(x)e^y & \sin(x)e^y \\ 0 & 3y^2 \end{bmatrix} \in \mathbb{R}^{3 \times 2}$$

```
from jax import vjp

def f(x):
    return jnp.array([x[0] + 2 * x[1], jnp.sin(x[0]) * jnp.exp(x[1]), x[1]**3])

def Jf_manual(x):
    return jnp.array([
        [1, 2],
        [jnp.cos(x[0]) * jnp.exp(x[1]), jnp.sin(x[0]) * jnp.exp(x[1])],
        [0, 3 * x[1] **2]
    ])

# compute the vjp manually
def vjp_manual(x, v):
    return v @ Jf_manual(x)

x = jnp.array([1., 2.])
v = jnp.array([2., 3., 4.])

print(f"manually computed jvp: {vjp_manual(x, v)}")

# compute the jvp using jax.jvp():

func_value, vjp_JAX = vjp(f, x) # vjp signature is a bit different from jvp
print(f"JAX computed jvp: {vjp_JAX(v)}")
```

✓ 0.2s

Python

vjp() example

$$f(x, y) = (x + 2y, \sin(x)e^y, y^3), f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$Jf(x, y) = \begin{bmatrix} 1 & 2 \\ \cos(x)e^y & \sin(x)e^y \\ 0 & 3y^2 \end{bmatrix} \in \mathbb{R}^{3 \times 2}$$

```
from jax import vjp

def f(x):
    return jnp.array([x[0] + 2 * x[1], jnp.sin(x[0]) * jnp.exp(x[1]), x[1]**3])

def Jf_manual(x):
    return jnp.array([
        [1, 2],
        [jnp.cos(x[0]) * jnp.exp(x[1]), jnp.sin(x[0]) * jnp.exp(x[1])],
        [0, 3 * x[1] **2]
    ])

# compute the vjp manually
def vjp_manual(x, v):
    return v @ Jf_manual(x)

x = jnp.array([1., 2.])
v = jnp.array([2., 3., 4.])

print(f"manually computed jvp: {vjp_manual(x, v)}")

# compute the jvp using jax.jvp():

func_value, vjp_JAX = vjp(f, x) # vjp signature is a bit different from jvp
print(f"JAX computed jvp: {vjp_JAX(v)}")
```

✓ 0.2s

Python

manually computed jvp: [13.976972 70.65303]

JAX computed jvp: (DeviceArray([13.976972, 70.65303], dtype=float32),)

Computing a full Jacobian matrix using `vjp()`

Post-multiplying a matrix with the i th standard basis vector reveals the i th column of the matrix, so can construct the full $m \times n$ Jacobian matrix with m VJPs:

```
def Jf_VJPs(x): # compute the full Jacobian at x using VJPs
    Jf = np.zeros(shape = (3, 2))
    for i in range(3):
        # compute the ith basis vector
        ei = np.zeros(shape = 3)
        ei[i] = 1

        # set the ith column equal to JVP with v = ej
        Jf[i] = vjp(f, x)[1](ei)[0]
    return Jf

x = jnp.array([2., 3.])

print(f"Jf with VJPs: \n {Jf_VJPs(x)}")
print(f"Jf manual: \n {Jf_manual(x)}")
```

✓ 0.1s

Python

Computing a full Jacobian matrix using `vjp()`

Post-multiplying a matrix with the i th standard basis vector reveals the i th column of the matrix, so can construct the full $m \times n$ Jacobian matrix with m VJPs:

```
def Jf_VJPs(x): # compute the full Jacobian at x using VJPs
    Jf = np.zeros(shape = (3, 2))
    for i in range(3):
        # compute the ith basis vector
        ei = np.zeros(shape = 3)
        ei[i] = 1

        # set the ith column equal to JVP with v = ej
        Jf[i] = vjp(f, x)[1](ei)[0]
    return Jf
```

```
x = jnp.array([2., 3.])
```

```
print(f"Jf with VJPs: \n {Jf_VJPs(x)}")
print(f"Jf manual: \n {Jf_manual(x)}")
```

✓ 0.1s

Python

Jf with VJPs:

```
[[ 1.          2.         ]
 [-8.35853291 18.26372719]
 [ 0.          27.         ]]
```

Jf manual:

```
[[ 1.          2.         ]
 [-8.358533 18.263727]
 [ 0.          27.         ]]
```


Reverse-mode differentiation: `jacrev()`

```
from jax import jacrev
```

```
print(f"Jf with jacrev(): \n {jacrev(f)(x)}")
```

✓ 0.3s

Python

```
Jf with jacrev():
```

```
[[ 1.          2.          ]
```

```
[-8.358533 18.263727]
```

```
[ 0.          27.          ]]
```

Reverse-mode differentiation: `jacrev()`

```
from jax import jacrev
```

```
print(f"Jf with jacrev(): \n {jacrev(f)(x)}")
```

✓ 0.3s

Python

```
Jf with jacrev():
```

```
[[ 1.          2.          ]
```

```
[-8.358533 18.263727]
```

```
[ 0.          27.          ]]
```

...better for short and
fat Jacobians

Forward vs. Reverse Mode Performance Comparison


$$f(x) = \left\| \frac{e^x - x^2}{\tan x} \right\|, x \in \mathbb{R}^n$$

$$g(t) = t^{\sin(1,t,t^2,\dots,t^{n-1})}, t \in \mathbb{R}$$

Forward vs. Reverse Mode Performance Comparison

$$f(x) = \left\| \frac{e^x - x^2}{\tan x} \right\|, x \in \mathbb{R}^n$$

$$g(t) = t^{\sin(1,t,t^2,\dots,t^{n-1})}, t \in \mathbb{R}$$



Performance

JAX comes with JIT compilation

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph
- Jitting (just-in-time compilation) compiles the computation graph and makes things run much faster

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph
- Jitting (just-in-time compilation) compiles the computation graph and makes things run much faster
 - Jitting traces through the entire computation (through function calls), so only need to jit outermost function

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph
- Jitting (just-in-time compilation) compiles the computation graph and makes things run much faster
 - Jitting traces through the entire computation (through function calls), so only need to jit outermost function
- Two limitations

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph
- Jitting (just-in-time compilation) compiles the computation graph and makes things run much faster
 - Jitting traces through the entire computation (through function calls), so only need to jit outermost function
- Two limitations
 - Control flow must be agnostic to values of traced variables

JAX comes with JIT compilation

- Automatic differentiation involves *tracing* a sequence of arithmetic computations to build up the computation graph
- Jitting (just-in-time compilation) compiles the computation graph and makes things run much faster
 - Jitting traces through the entire computation (through function calls), so only need to jit outermost function
- Two limitations
 - Control flow must be agnostic to value of traced variables
 - Requires array shapes of functions to be known ahead of time

JIT example, simple function

$$f(x) = \sigma \left(\left\| \hat{I} - I \right\|_2^2 \right), I \in \mathbb{R}^{128 \times 128}$$

JIT example, simple function

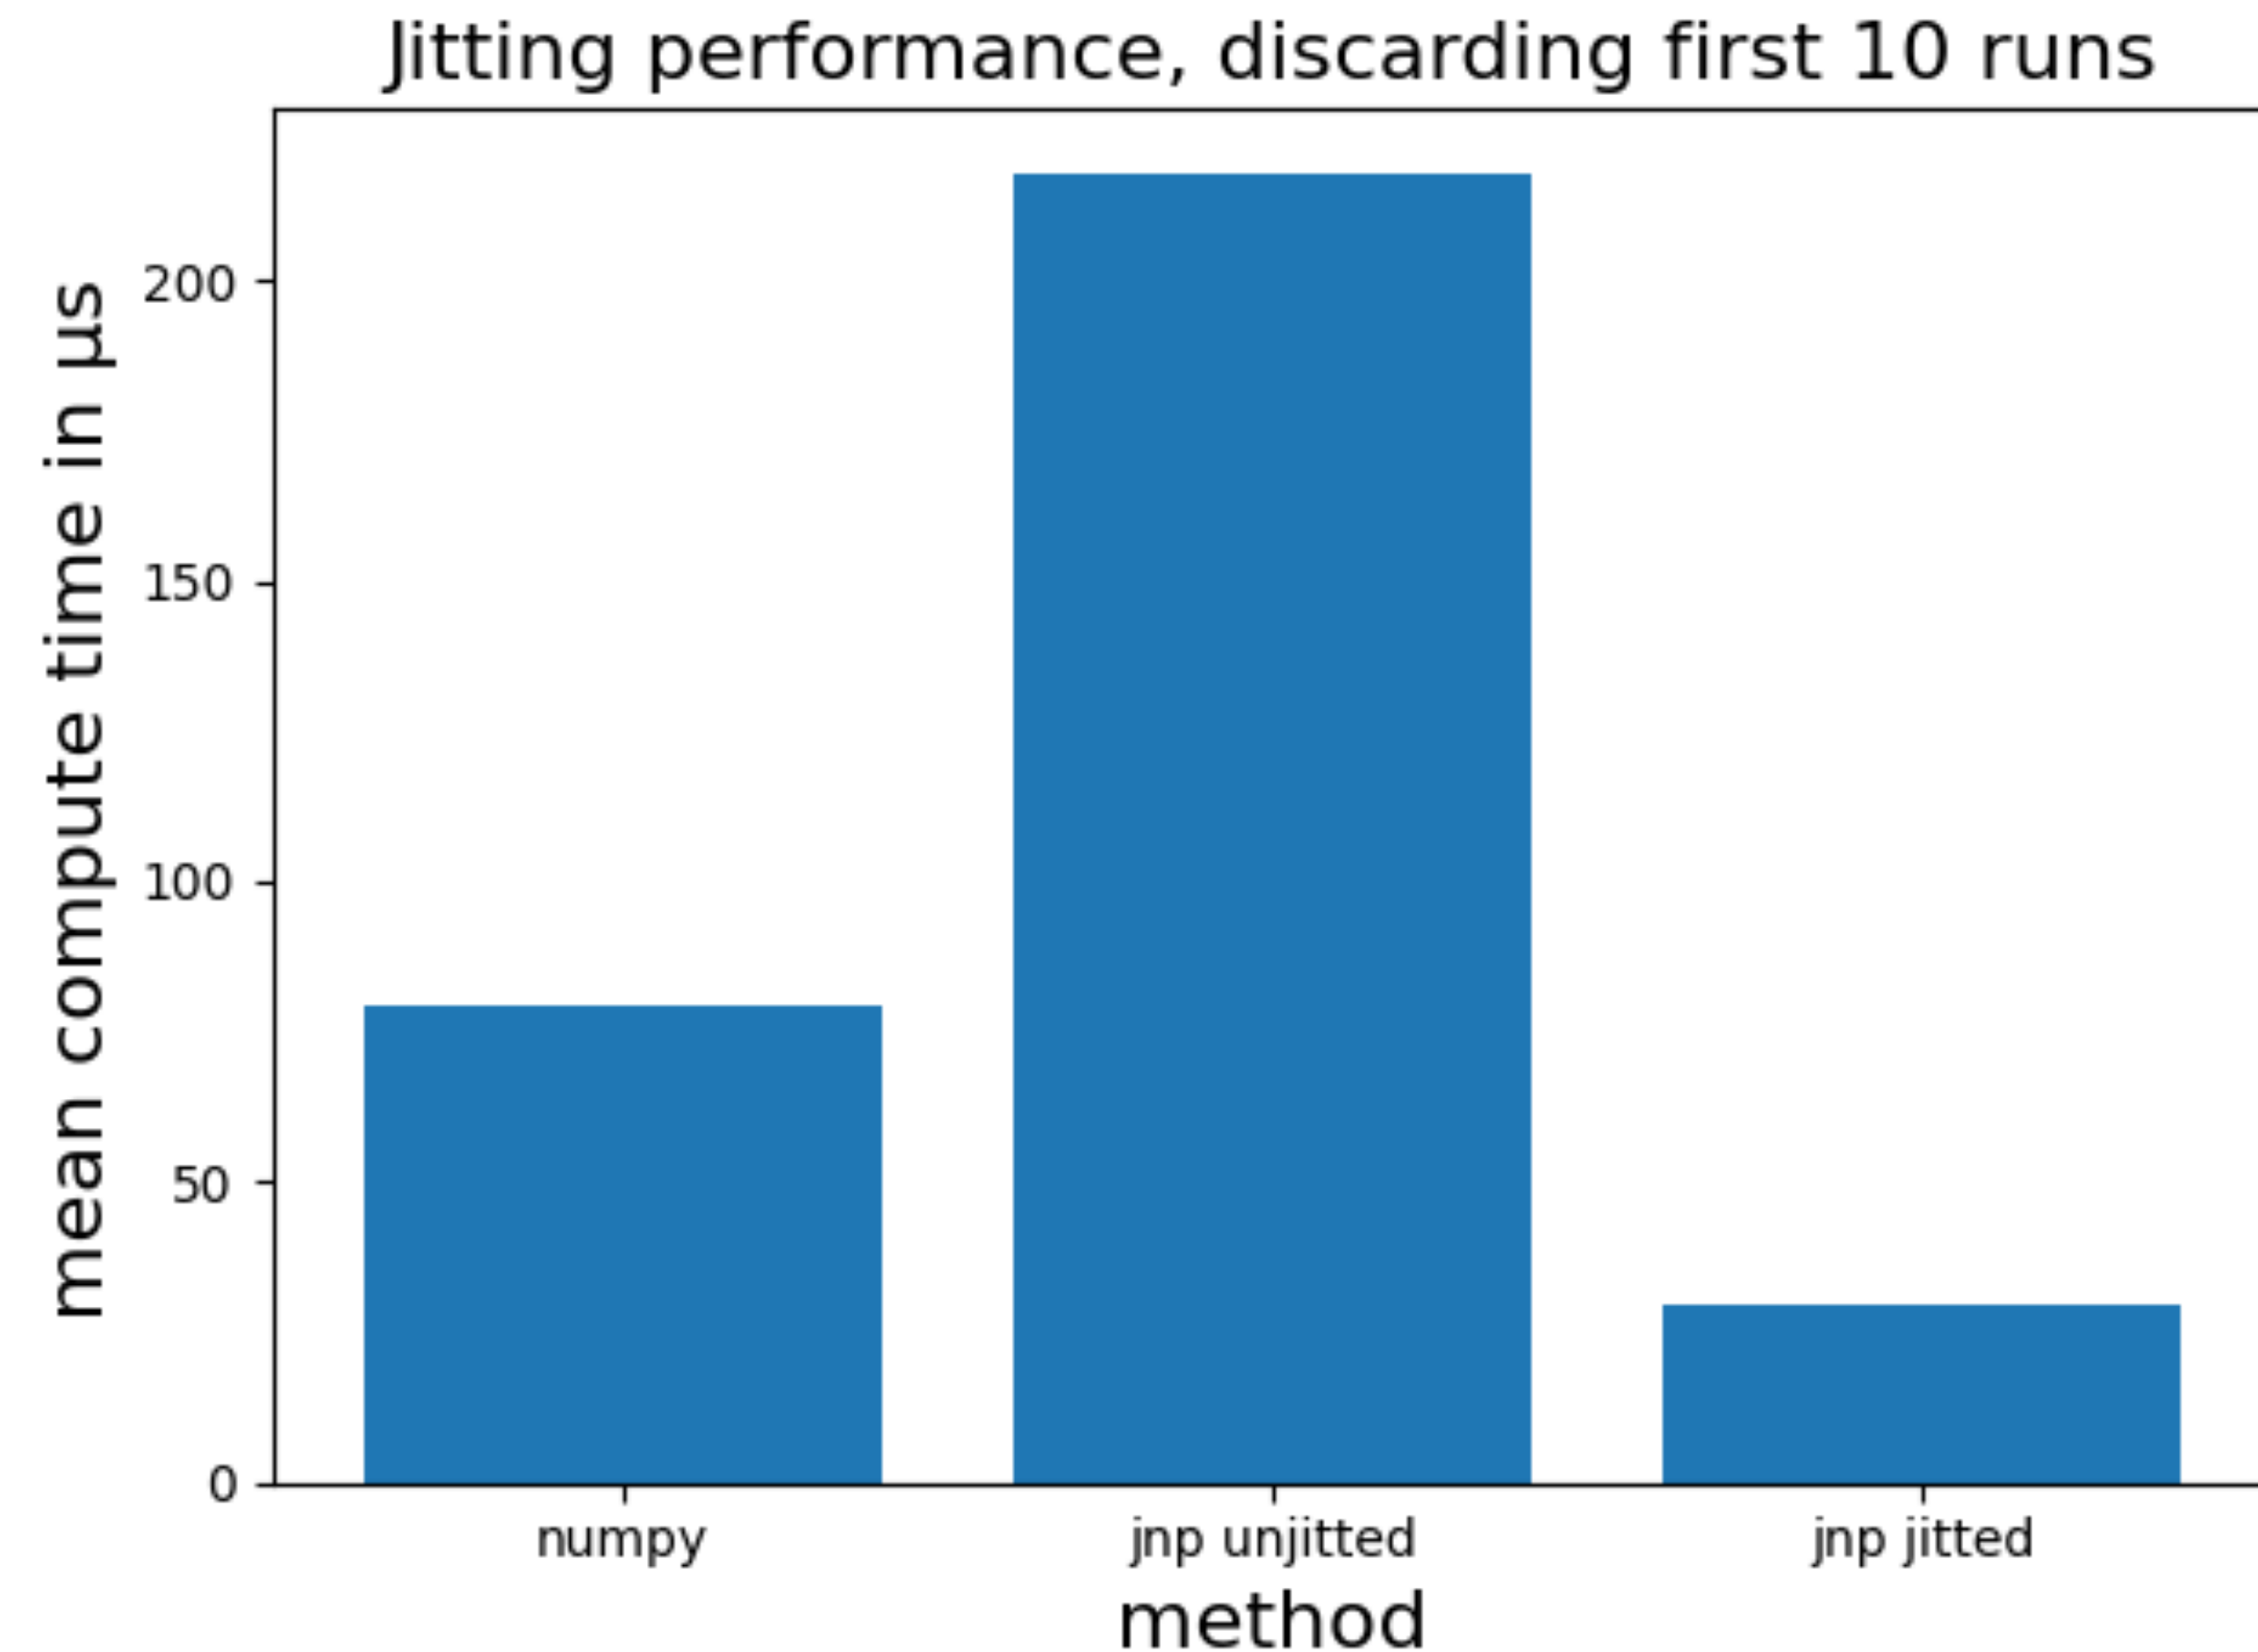
$$f(x) = \sigma \left(\left\| \hat{I} - I \right\|_2^2 \right), I \in \mathbb{R}^{128 \times 128}$$

```
1  def f_np(Ihat, I):
2      diff = np.linalg.norm(Ihat - I) **2 / 2
3      return 1 / (1 + np.exp(-diff))
4
5
6  def f_jnp(Ihat, I):
7      diff = jnp.linalg.norm(Ihat - I) **2 / 2
8      return 1 / (1 + jnp.exp(-diff))
9
10 f_jnp_jitted = jit(f_jnp)
```

JIT example, simple function

$$f(x) = \sigma \left(\left\| \hat{I} - I \right\|_2^2 \right), I \in \mathbb{R}^{128 \times 128}$$

```
1 def f_np(Ihat, I):
2     diff = np.linalg.norm(Ihat - I) **2 / 2
3     return 1 / (1 + np.exp(-diff))
4
5
6 def f_jnp(Ihat, I):
7     diff = jnp.linalg.norm(Ihat - I) **2 / 2
8     return 1 / (1 + jnp.exp(-diff))
9
10 f_jnp_jitted = jit(f_jnp)
```



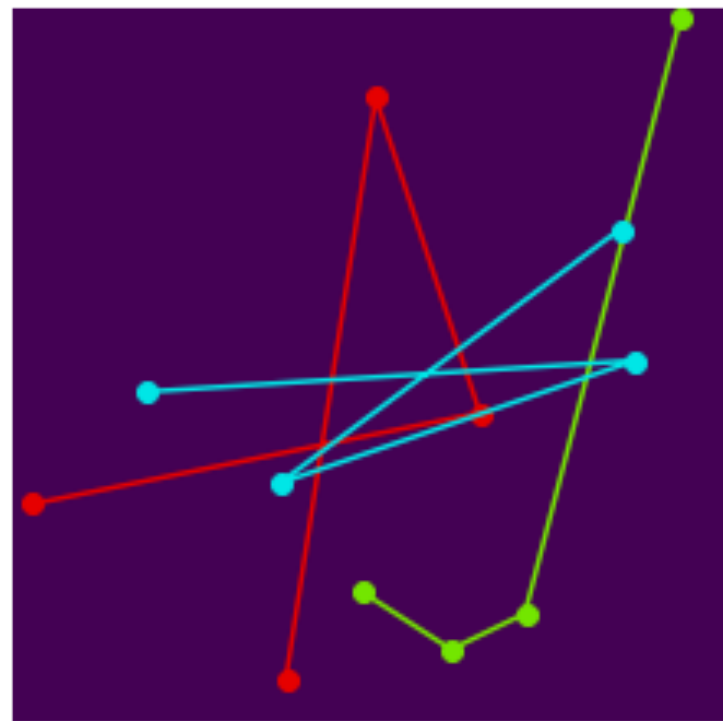
JIT example, Bezier rendering function

$$\mathbf{B}(t) = (1 - t)^3 \mathbf{P}_0 + 3(1 - t)^2 t \mathbf{P}_1 + 3(1 - t) t^2 \mathbf{P}_2 + t^3 \mathbf{P}_3, \quad \begin{array}{l} t \in [0,1] \\ \mathbf{P}_i \in \mathbb{R}^2, \text{ control points} \end{array}$$


JIT example, Bezier rendering function

$$\mathbf{B}(t) = (1 - t)^3 \mathbf{P}_0 + 3(1 - t)^2 t \mathbf{P}_1 + 3(1 - t) t^2 \mathbf{P}_2 + t^3 \mathbf{P}_3, \quad t \in [0, 1]$$

$\mathbf{P}_i \in \mathbb{R}^2$, control points




Control Points





Rendered Image

JIT example, Bezier rendering function


$$\mathbf{B}(t) = (1 - t)^3 \mathbf{P}_0 + 3(1 - t)^2 t \mathbf{P}_1 + 3(1 - t) t^2 \mathbf{P}_2 + t^3 \mathbf{P}_3, \quad \begin{array}{l} t \in [0,1] \\ \mathbf{P}_i \in \mathbb{R}^2, \text{ control points} \end{array}$$



Control Points



Rendered Image



Control flow can't depend on values of inputs

```
from jax import jit

def add_one_or_two(x, add_one):
    if add_one: # control flow depends on the argument add_one
        return x + 1
    else:
        return x + 2

add_one_or_two_jitted = jit(add_one_or_two)
add_one_or_two_jitted(1.0, True)
```

⊗ 1.5s

Python

Control flow can't depend on values of inputs

```
from jax import jit

def add_one_or_two(x, add_one):
    if add_one: # control flow depends on the argument add_one
        return x + 1
    else:
        return x + 2

add_one_or_two_jitted = jit(add_one_or_two)
add_one_or_two_jitted(1.0, True)
```

⊗ 1.5s

Python

```
-----
ConcretizationError                                Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 41' in <cell line: 10>()
      7     return x + 2
      9 add_one_or_two_jitted = jit(add_one_or_two)
--> 10 add_one_or_two_jitted(1.0, True)
```

[... skipping hidden 14 frame]

```
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 41' in add_one_or_two(x, add_one)
      3 def add_one_or_two(x, add_one):
--> 4     if add_one: # control flow depends on the argument add_one
      5         return x + 1
      6     else:
```

[... skipping hidden 1 frame]

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/core.py:1123, in concretization_function_error.<locals>.error(self, arg)
    1122 def error(self, arg):
-> 1123     raise ConcretizationError(arg, fname_context)
```

ConcretizationTypeError: Abstract tracer value encountered where concrete value is expected: Traced<ShapedArray(bool[], weak_type=True)>with<DynamicJaxprTrace(level=0/1)>
The problem arose with the 'bool' function.
While tracing the function add_one_or_two at /tmp/ipykernel_1263005/1343836691.py:3 for jit, this concrete value was not available in Python because it depends on the value of the argument 'add_one'.

See <https://jax.readthedocs.io/en/latest/errors.html#jax.errors.ConcretizationTypeError>

Use static_argnums

```
def add_one_or_two(x, add_one):  
    if add_one: # control flow depends on the argument add_one  
        return x + 1  
    else:  
        return x + 2  
  
add_one_or_two_jitted = jit(add_one_or_two, static_argnums = 1)  
add_one_or_two_jitted(1.0, True)
```

✓ 0.1s

Python

DeviceArray(2., dtype=float32, weak_type=True)

Array shapes can't depend on values of inputs

```
def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
    bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
    ab = np.sum(bs)

    return ab
```

```
stupid_mult_jitted = jit(stupid_mult)
stupid_mult_jitted(3, 4)
```

⊗ 0.9s

Python

Array shapes can't depend on values of inputs

```
def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
    bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
    ab = np.sum(bs)

    return ab
```

```
stupid_mult_jitted = jit(stupid_mult)
stupid_mult_jitted(3, 4)
```

⊗ 0.9s

Python

```
-----
TracerIntegerConversionError      Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in <cell line: 8>()
      5     return ab
      7     stupid_mult_jitted = jit(stupid_mult)
---->  8     stupid_mult_jitted(3, 4)
```

[... skipping hidden 14 frame]

```
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in stupid_mult(a, b)
      1 def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
---->  2     bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
      3     ab = np.sum(bs)
      5     return ab
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/numpy/core/numeric.py:204, in ones(shape, dtype, order, like)
      201 if like is not None:
      202     return _ones_with_like(shape, dtype=dtype, order=order, like=like)
-->  204 a = empty(shape, dtype, order)
      205 multiarray.copyto(a, 1, casting='unsafe')
      206 return a
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/core.py:519, in Tracer.__index__(self)
      518 def __index__(self):
-->  519     raise TracerIntegerConversionError(self)
```

TracerIntegerConversionError: The __index__() method was called on the JAX Tracer object Traced<ShapedArray(int32[], weak_type=True)>with<DynamicJaxprTrace(level=0/1)>
See <https://jax.readthedocs.io/en/latest/errors.html#jax.errors.TracerIntegerConversionError>

Array shapes can't depend on values of inputs

```
def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
    bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
    ab = np.sum(bs)

    return ab
```

```
stupid_mult_jitted = jit(stupid_mult)
stupid_mult_jitted(3, 4)
```

⊗ 0.9s

Python

```
-----
TracerIntegerConversionError      Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in <cell line: 8>()
      5     return ab
      7     stupid_mult_jitted = jit(stupid_mult)
---->  8     stupid_mult_jitted(3, 4)
```

[... skipping hidden 14 frame]

```
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in stupid_mult(a, b)
      1 def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
---->  2     bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
      3     ab = np.sum(bs)
      5     return ab
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/numpy/core/numeric.py:204, in ones(shape, dtype, order, like)
    201 if like is not None:
    202     return _ones_with_like(shape, dtype=dtype, order=order, like=like)
-->  204 a = empty(shape, dtype, order)
    205 multiarray.copyto(a, 1, casting='unsafe')
    206 return a
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/core.py:519, in Tracer.__index__(self)
    518 def __index__(self):
-->  519     raise TracerIntegerConversionError(self)
```

TracerIntegerConversionError: The __index__() method was called on the JAX Tracer object Traced<ShapedArray(int32[], weak_type=True)>with<DynamicJaxprTrace(level=0/1)>
See <https://jax.readthedocs.io/en/latest/errors.html#jax.errors.TracerIntegerConversionError>

```
stupid_mult_jitted = jit(stupid_mult, static_argnums = 0)
stupid_mult_jitted(3, 4)
```

✓ 0.1s

Python

Array shapes can't depend on values of inputs

```
def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
    bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
    ab = np.sum(bs)

    return ab
```

```
stupid_mult_jitted = jit(stupid_mult)
stupid_mult_jitted(3, 4)
```

⊗ 0.9s

Python

```
-----
TracerIntegerConversionError      Traceback (most recent call last)
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in <cell line: 8>()
      5     return ab
      7 stupid_mult_jitted = jit(stupid_mult)
---->  8 stupid_mult_jitted(3, 4)
```

[... skipping hidden 14 frame]

```
/data/philiplo125/jax-tutorial/tutorial.ipynb Cell 45' in stupid_mult(a, b)
      1 def stupid_mult(a, b): # multiply two integers a and b in a very dumb way
---->  2     bs = b * np.ones(shape = (a,)) # the shape of bs depends on a!
      3     ab = np.sum(bs)
      5     return ab
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/numpy/core/numeric.py:204, in ones(shape, dtype, order, like)
    201 if like is not None:
    202     return _ones_with_like(shape, dtype=dtype, order=order, like=like)
-->  204 a = empty(shape, dtype, order)
    205 multiarray.copyto(a, 1, casting='unsafe')
    206 return a
```

```
File ~/anaconda3/envs/jax-tutorial/lib/python3.9/site-packages/jax/core.py:519, in Tracer.__index__(self)
    518 def __index__(self):
-->  519     raise TracerIntegerConversionError(self)
```

TracerIntegerConversionError: The __index__() method was called on the JAX Tracer object Traced<ShapedArray(int32[], weak_type=True)>with<DynamicJaxprTrace(level=0/1)>
See <https://jax.readthedocs.io/en/latest/errors.html#jax.errors.TracerIntegerConversionError>

```
stupid_mult_jitted = jit(stupid_mult, static_argnums = 0)
stupid_mult_jitted(3, 4)
```

✓ 0.1s

Python

DeviceArray(12., dtype=float32)

Automatic vectorization with `vmap()`

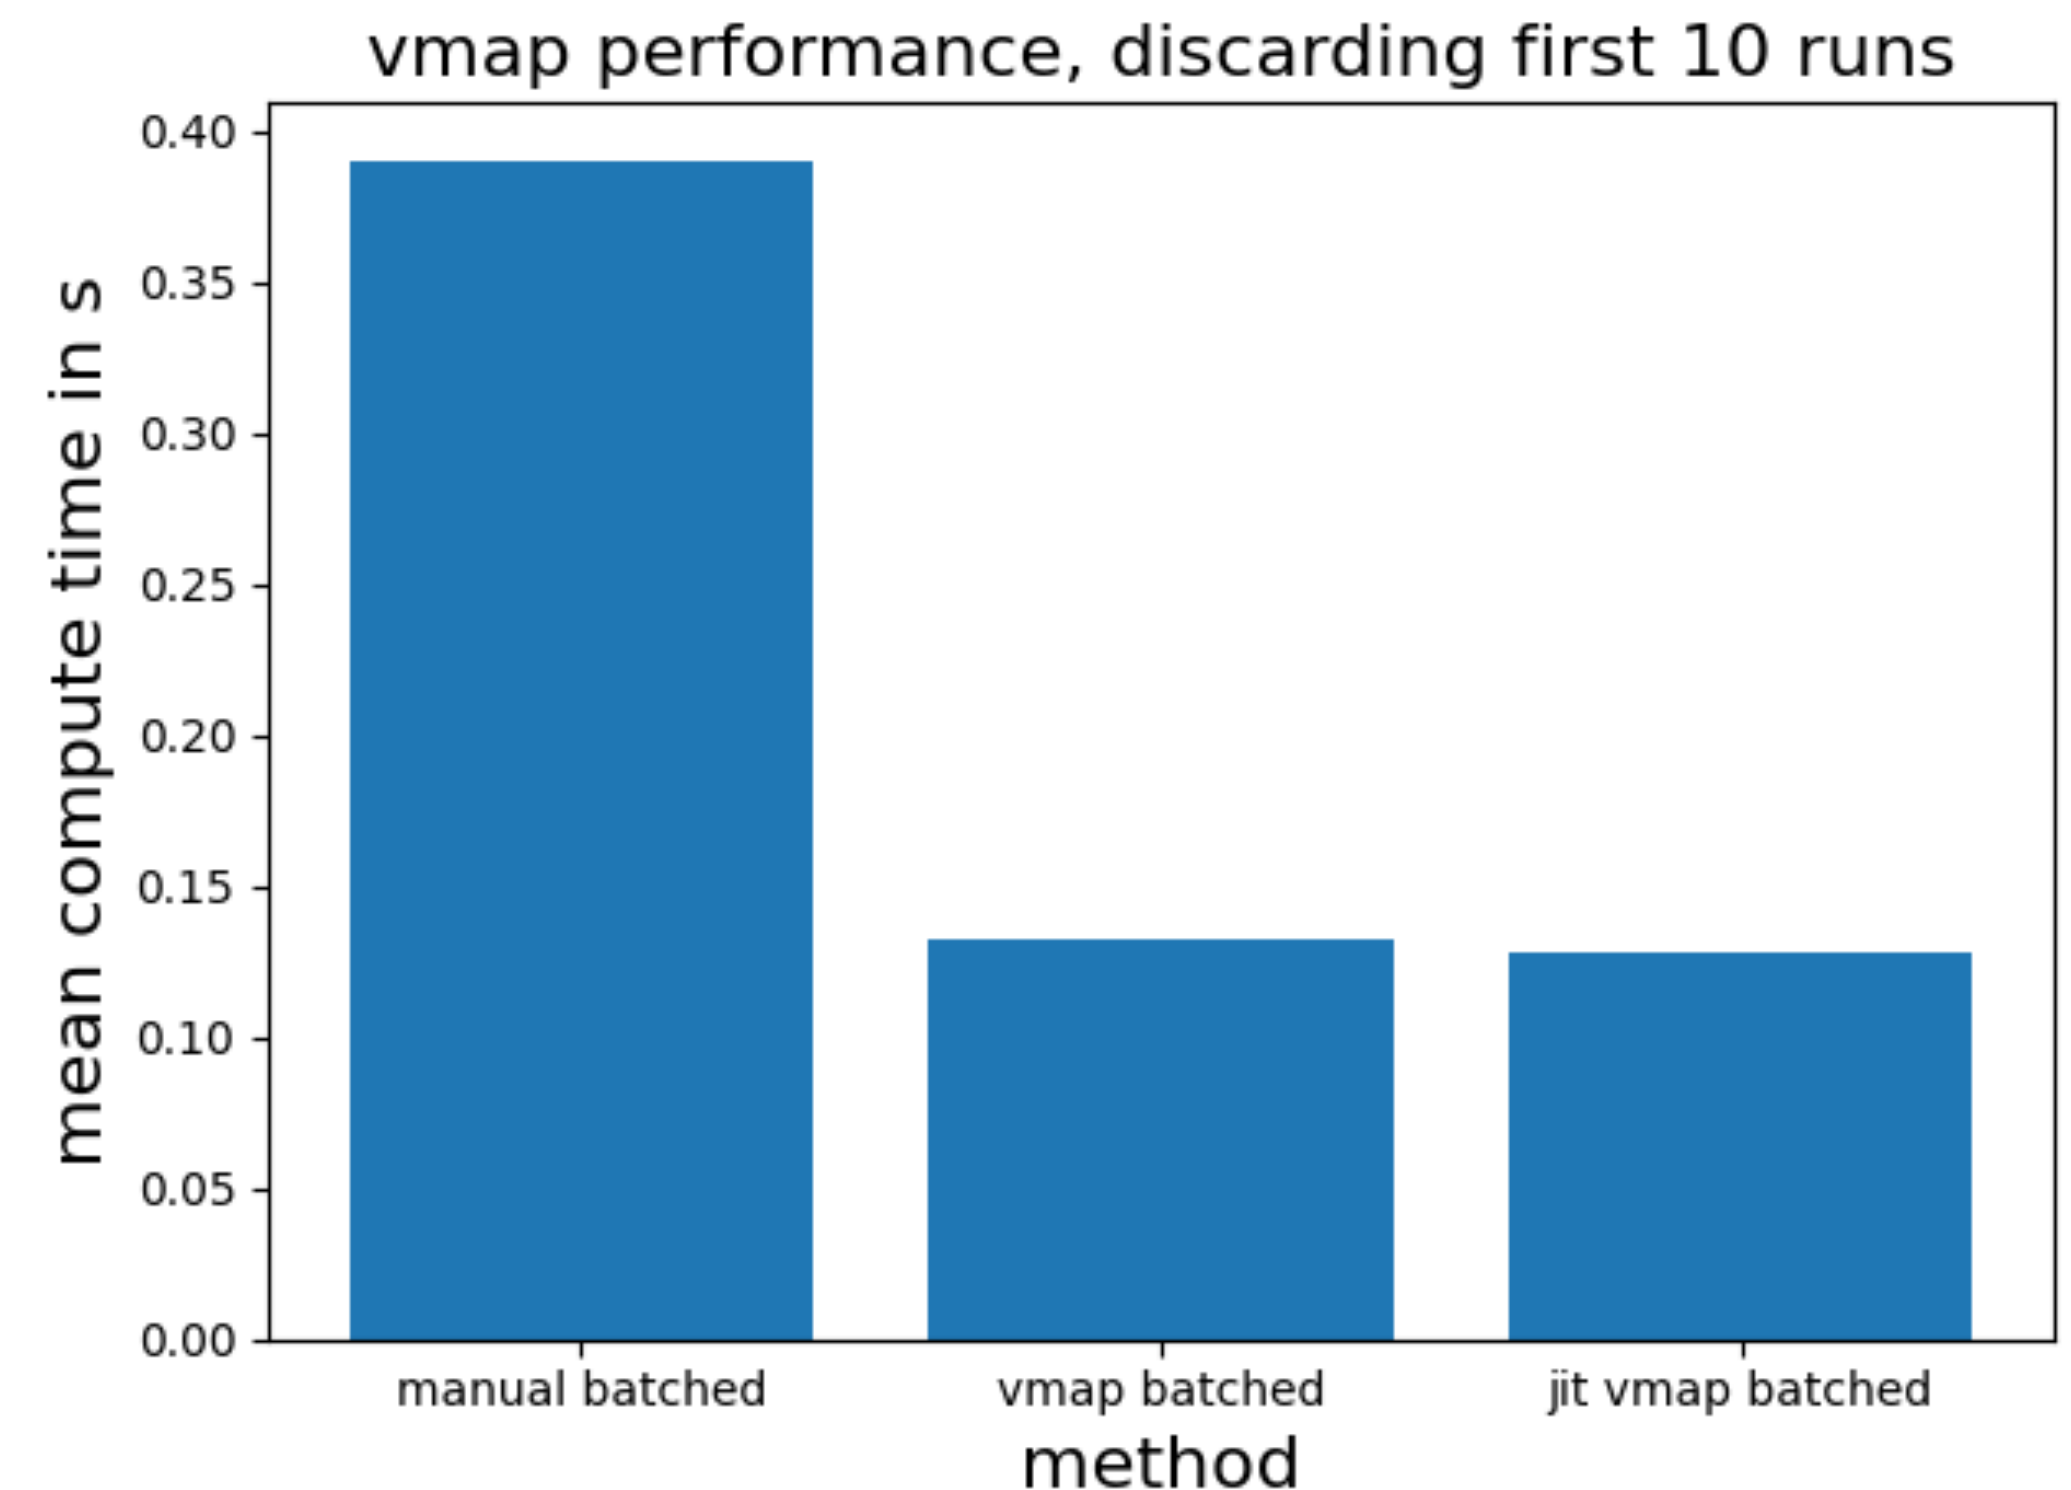
Compute Frobenius norm of one hundred 1000 x 1000 matrices

```
1  def frob_norm(A):  
2      '''  
3      Compute the Frobenius norm of a matrix A  
4      '''  
5      return jnp.linalg.norm(A)  
6  
7  def manual_batched_frob_norm(As):  
8      n = As.shape[0]  
9      norms = jnp.zeros(shape = As.shape[0])  
10     for i in range(n):  
11         norms = norms.at[i].set(frob_norm(As[i]))  
12     return norms  
13  
14  vmap_batched_frob_norm = vmap(frob_norm)  
15  
16  jitted_vmap_batched_frob_norm = jit(vmap_batched_frob_norm)
```

Automatic vectorization with `vmap()`

Compute Frobenius norm of one hundred 1000 x 1000 matrices

```
1  def frob_norm(A):  
2      '''  
3      Compute the Frobenius norm of a matrix A  
4      '''  
5      return jnp.linalg.norm(A)  
6  
7  def manual_batched_frob_norm(As):  
8      n = As.shape[0]  
9      norms = jnp.zeros(shape = As.shape[0])  
10     for i in range(n):  
11         norms = norms.at[i].set(frob_norm(As[i]))  
12     return norms  
13  
14  vmap_batched_frob_norm = vmap(frob_norm)  
15  
16  jitted_vmap_batched_frob_norm = jit(vmap_batched_frob_norm)
```



Other things you can do with JAX

Other things you can do with JAX

- Can arbitrarily compose forward and reverse mode differentiation

Other things you can do with JAX

- Can arbitrarily compose forward and reverse mode differentiation
- Can compute derivatives for functions between tensors of arbitrary shape

Other things you can do with JAX

- Can arbitrarily compose forward and reverse mode differentiation
- Can compute derivatives for functions between tensors of arbitrary shape
- Deep learning library: flax

Other things you can do with JAX

- Can arbitrarily compose forward and reverse mode differentiation
- Can compute derivatives for functions between tensors of arbitrary shape
- Deep learning library: flax
- Can run transparently on GPUs and TPUs, compiled with XLA

Other things you can do with JAX

- Can arbitrarily compose forward and reverse mode differentiation
- Can compute derivatives for functions between tensors of arbitrary shape
- Deep learning library: flax
- Can run transparently on GPUs and TPUs, compiled with XLA
- Repo: <https://github.com/PhillipLo/jax-tutorial>